

SENTINEL 2026

Must Have / Must Not Have · Quick Checklist

MUST HAVE

- Defense makes one of 4 decisions: ALLOW / BLOCK / ESCALATE / REWRITE.
- Decisions use agent state + candidate action + provenance + policy + observed content.
- Track provenance / trust for observations and memory.
- Observability shows action → decision → risk/reason → outcome.
- Demonstrate at least one benign task completing normally.
- Demonstrate a real attack reaching the defense and being handled.
- Test the published attack families and scenario library.
- Technical report: threat model, hypothesis, method, results, ablation, failures, Responsible AI.
- Include at least one ablation study.
- Document limitations, false positives, observed data, and human-in-the-loop cases.
- Declare external models and datasets.
- Submit defense + observability + 5–10 min video + GitHub repo + technical report.

MUST NOT HAVE

- No scenario-specific hard-coding, scenario IDs, filenames, or expected outcomes.
- No keyword/pattern matching as the core defense.
- No blocking everything just because content is suspicious or untrusted.
- Do not treat untrusted email, documents, logs, tool output, or memory as authority.
- Do not automatically promote untrusted memory into trusted policy.
- Do not accept fake/untrusted approval for sensitive financial actions.
- Do not hide failures or claim complete safety.
- Do not use a static log dump that fails to explain decisions.
- Do not attack real systems, steal credentials, escape the sandbox, persist, DoS, or use real personal data.

Key principle: Read untrusted content, but do not automatically trust it. Protect the action while keeping the agent useful.

Based on the SENTINEL Specification Book, IndabaX Tunisia 2026.